

# Parallel Programming Final Project

Team 8  
112062520 戴維恩  
111062698 戴樂為

# Table of contents



## Image Filters

Six image filters we implemented.



## Implementation

Code implementation explanation.



## Experiments

Experimental results on the optimized code.

#1

# Image Filters

- Gaussian Blur
- Emboss
- Erosion
- Dilation
- Wave
- Oil Painting



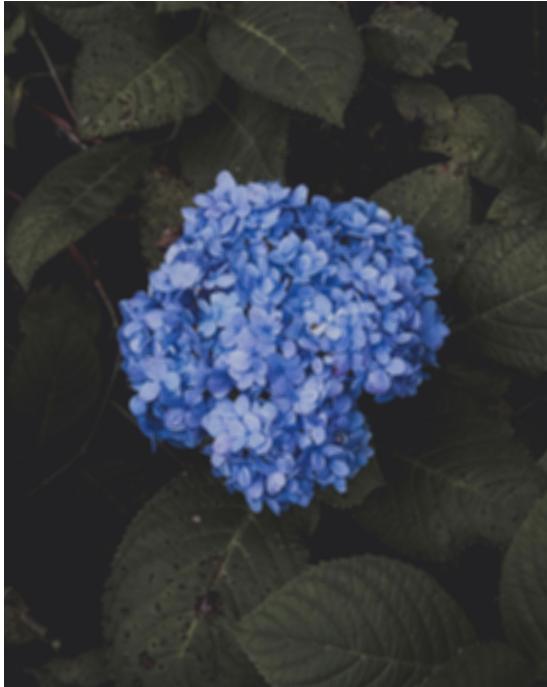
01

# Gaussian Blur

Parameters: Radius, Sigma



input



Radius = 10  
Sigma = 30



Radius = 30  
Sigma = 30



02

# Emboss

Parameters: Intensity



input



Intensity = 5



Intensity = 20



03

# Erosion

Parameters: Radius



input



Radius = 5

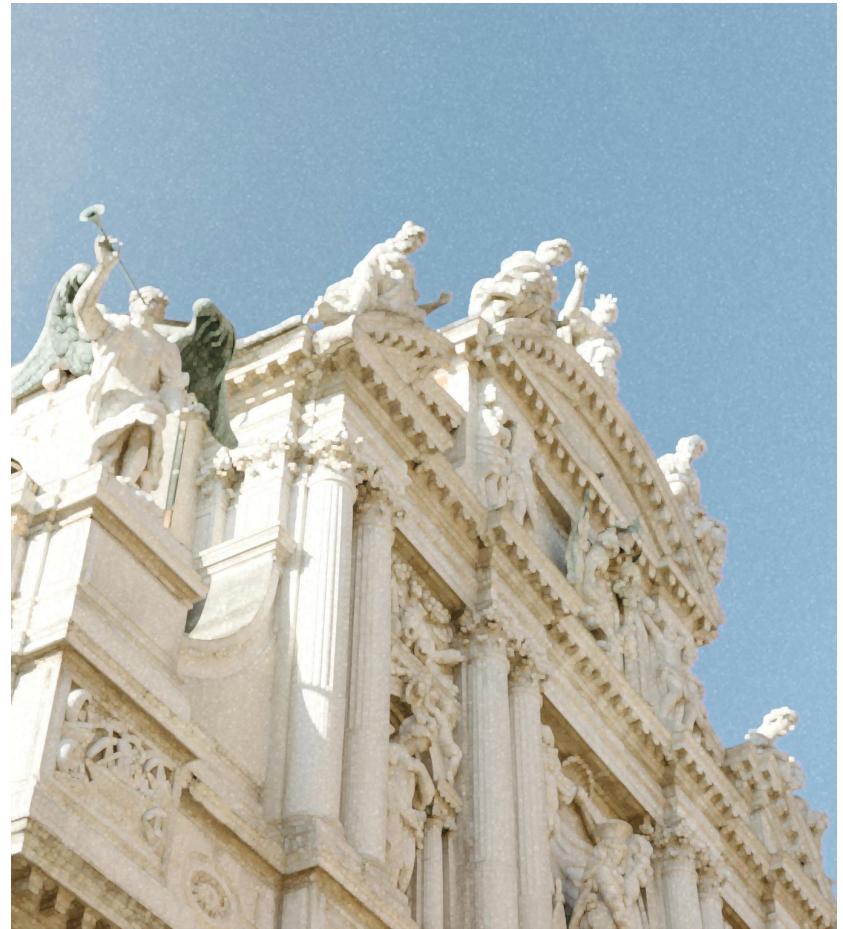


Radius = 10

# 04

# Dilation

Parameters: Radius





Input



Radius = 5



Radius = 10



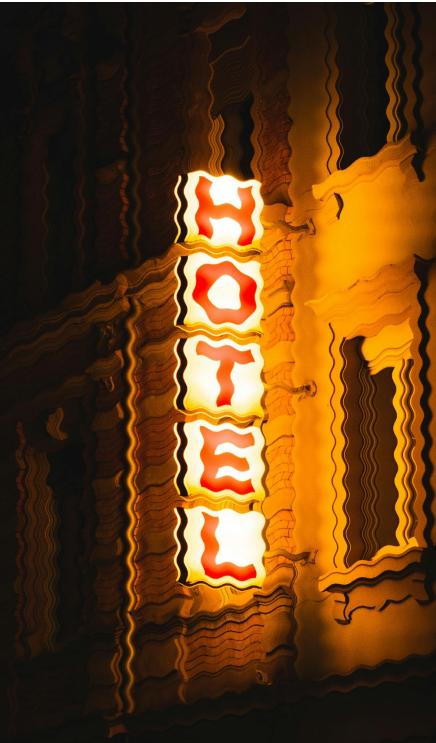
05

# Wave

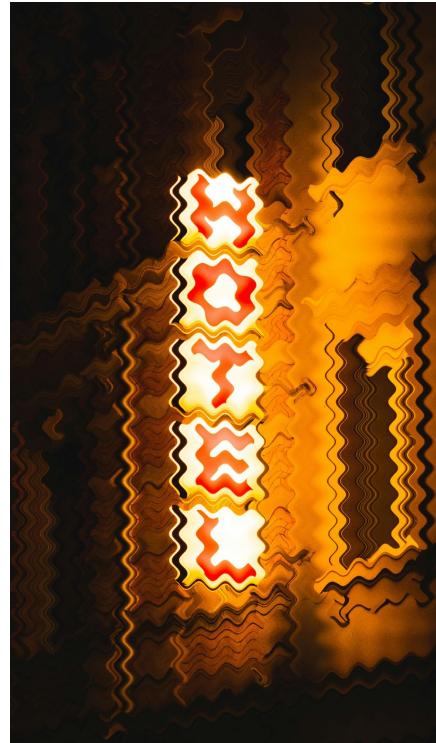
Parameters: Frequency, Amplitude



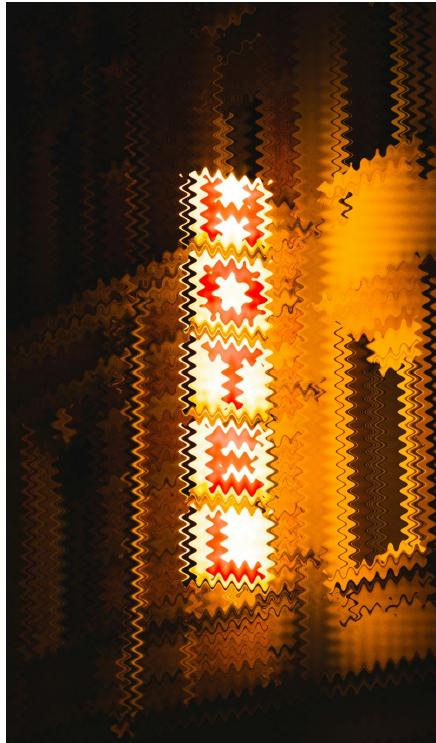
Input



Frequency = 0.05  
Amplitude = 10



Frequency = 0.05  
Amplitude = 20



Frequency = 0.1  
Amplitude = 20



# 06

# Oil Painting

Parameters: Radius



Input



Radius = 10



Radius = 15

#2

# Implementation

- Min/Max Operations
  - Erosion
  - Dilation
- Weighted Combination
  - Gaussian Blur
  - Emboss
- Special Operations
  - Wave (Pixel Displacement)
  - Oil Painting (Intensity-based operation)

# Setup

- Calculate Dimension

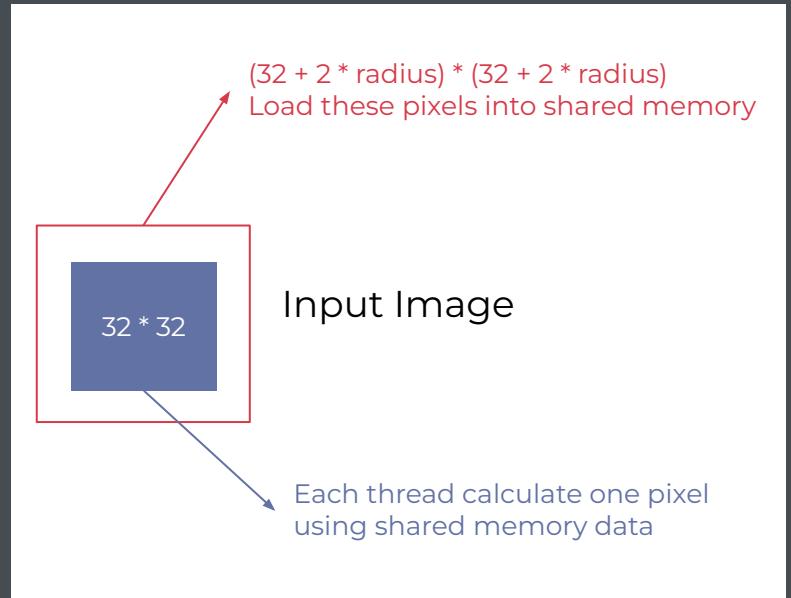
```
// code in kernel
const int tile_w = BLOCK_DIM_X + 2 * radius;
const int tile_h = BLOCK_DIM_Y + 2 * radius;
const int kernelSize = 2 * radius + 1;
```

- Coordinate Clamping

```
// clamp y to 0 ~ height-1
gy = max(0, min(gy, height - 1));
// clamp x to 0 ~ width-1
gx = max(0, min(gx, width - 1));
```

- Kernel Launch

```
// thread block & grid dimension
dim3 block(32, 32);
dim3 grid(
    (width + 32 - 1) / 32,
    (height + 32 - 1) / 32
);
```

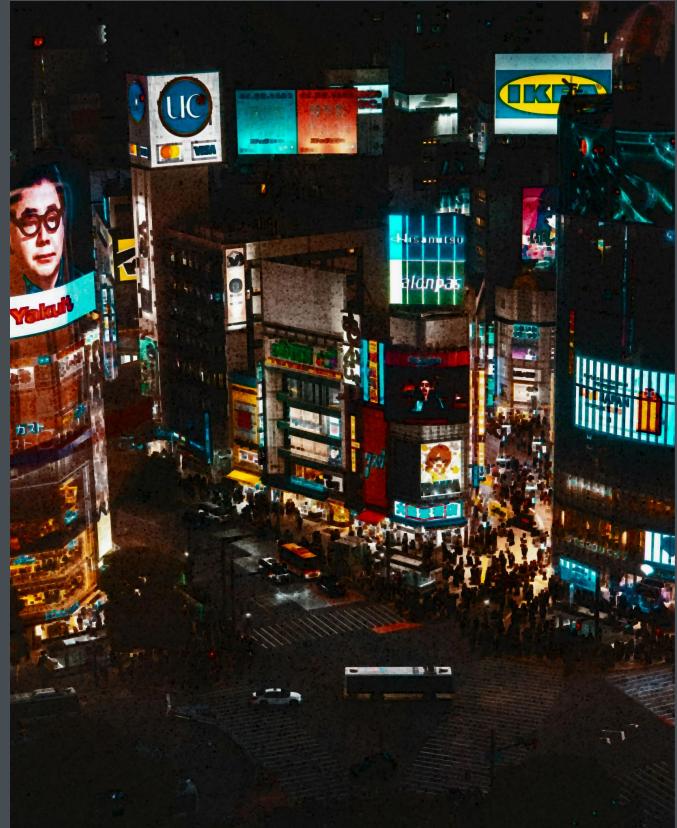


# Min/Max Operation - Erosion

- CUDA Code

```
unsigned char min_val = 255;
for(int ky = -radius; ky <= radius; ky++) {
    for(int kx = -radius; kx <= radius; kx++) {
        int py = min(max(y + ky, 0), height - 1);
        int px = min(max(x + kx, 0), width - 1);

        unsigned char val = src[(py * width + px) * channels + c];
        min_val = min(val, min_val);
    }
}
dst[idx] = min_val;
```



# Min/Max Operation - Dilation

- CUDA Code

```
unsigned char max_val = 0;  
for(int ky = -radius; ky <= radius; ky++) {  
    for(int kx = -radius; kx <= radius; kx++) {  
        int py = min(max(y + ky, 0), height - 1);  
        int px = min(max(x + kx, 0), width - 1);  
  
        unsigned char val = src[(py * width + px) * channels + c];  
        max_val = max(val, max_val);  
    }  
}  
dst[idx] = max_val;
```



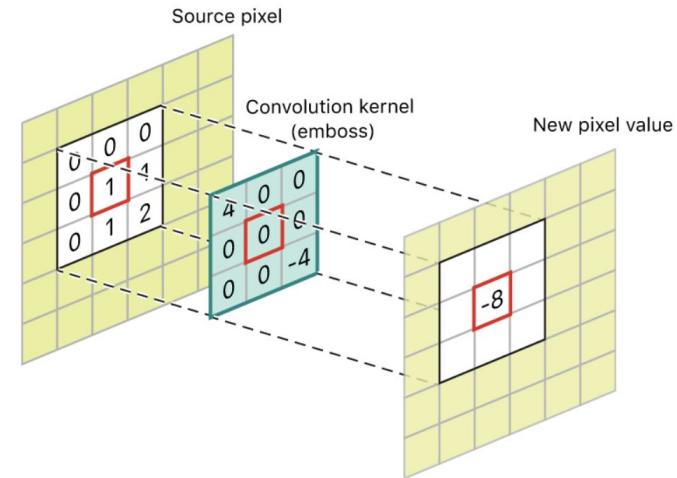
# Weighted Combination - Gaussian Blur

- Filter

```
// compute filter on CPU
for(int y = -radius; y <= radius; y++) {
    for(int x = -radius; x <= radius; x++) {
        float value = exp(-(x*x + y*y) / (2*sigma*sigma));
        kernel[(y+radius) * kernelSize + (x+radius)] = value;
        sum += value;
    }
}
```

- CUDA Code

```
for(int ky = -radius; ky <= radius; ky++) {
    for(int kx = -radius; kx <= radius; kx++)
        // ...calculate index
        val += src[src_idx] * kernel[kernel_idx];
    }
}
```



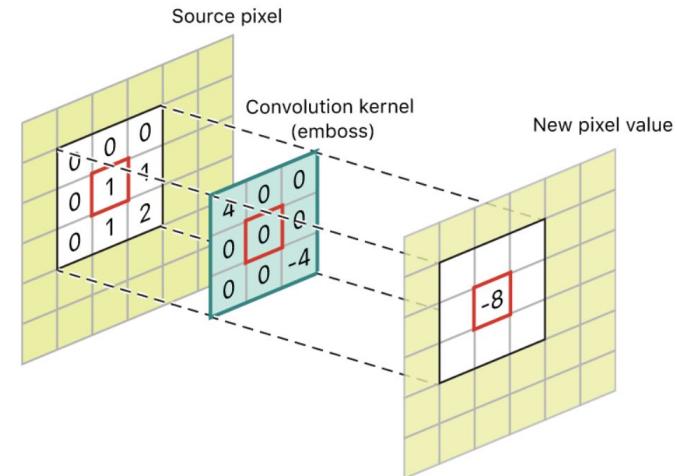
# Weighted Combination - Emboss

- Filter

```
// compute filter on CPU  
  
float side = -2.0f * intensity, corner = -1.0f * intensity;  
  
float center = 1.0f, opposite = 2.0f * intensity;  
  
float kernel[3][3] = {  
    {side, corner, 0},  
    {corner, center, intensity},  
    {0, intensity, opposite}  
};
```

- CUDA Code

```
for(int ky = -radius; ky <= radius; ky++) {  
    for(int kx = -radius; kx <= radius; kx++) {  
        // ...calculate index  
        val += src[src_idx] * kernel[kernel_idx];  
    }  
}
```



# Special Operation - Wave

- CUDA Code

```
// kernel code  
  
int x = blockIdx.x * blockDim.x + threadIdx.x;  
int y = blockIdx.y * blockDim.y + threadIdx.y;  
  
int sourceX = x + amplitudeX * sin(y * frequencyY);  
int sourceY = y + amplitudeY * sin(x * frequencyX);  
  
Output[y][x] = Input[sourceY][sourceX];
```



# Special Operation - Oil Painting

- CUDA Code

```
for (ky = -radius; ky <= radius; ky++) {  
    for (kx = -radius; kx <= radius; kx++) {  
        intensity = (img[py + ky][px + kx] * maxLevels) / 256;  
        count[intensity]++;  
        sum[intensity] += img[py + ky][px + kx];  
    }  
}  
  
// ... find the intensity level I with maximum count.  
Output[py][py] = sum[I] / count[I];
```

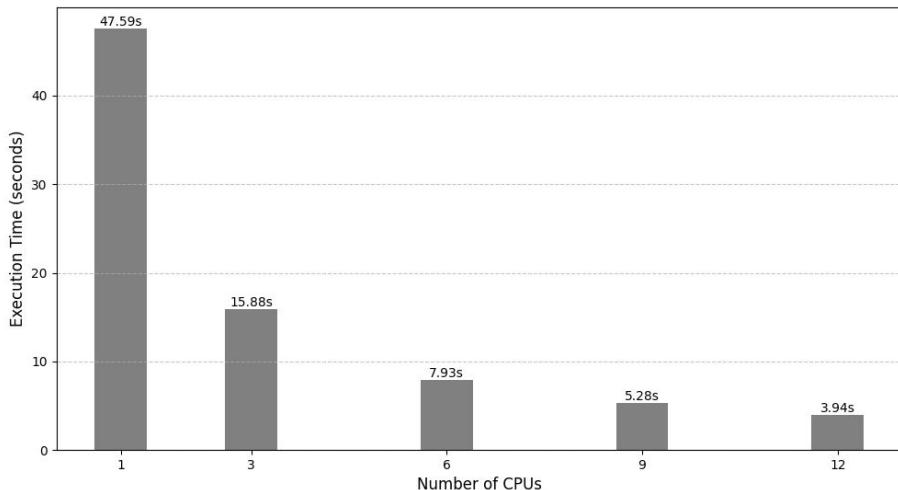


# #3 Experiments

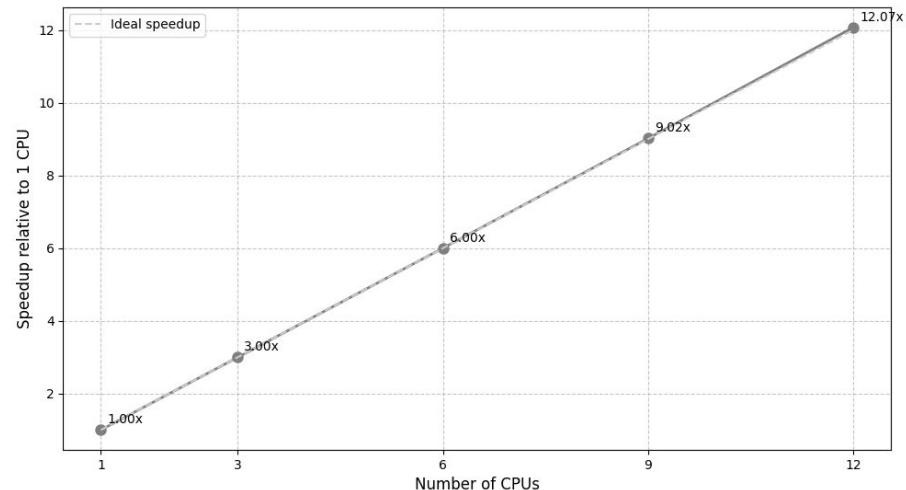
- OpenMP
  - Time profile & Speedup
- Cuda
  - Baseline vs Optimized
  - Blocking Factor
  - Time Distribution

# OpenMP: Time Profile & Speedup

Execution Time vs Number of CPUs



Performance Speedup vs Number of CPUs



# CUDA: nvprof

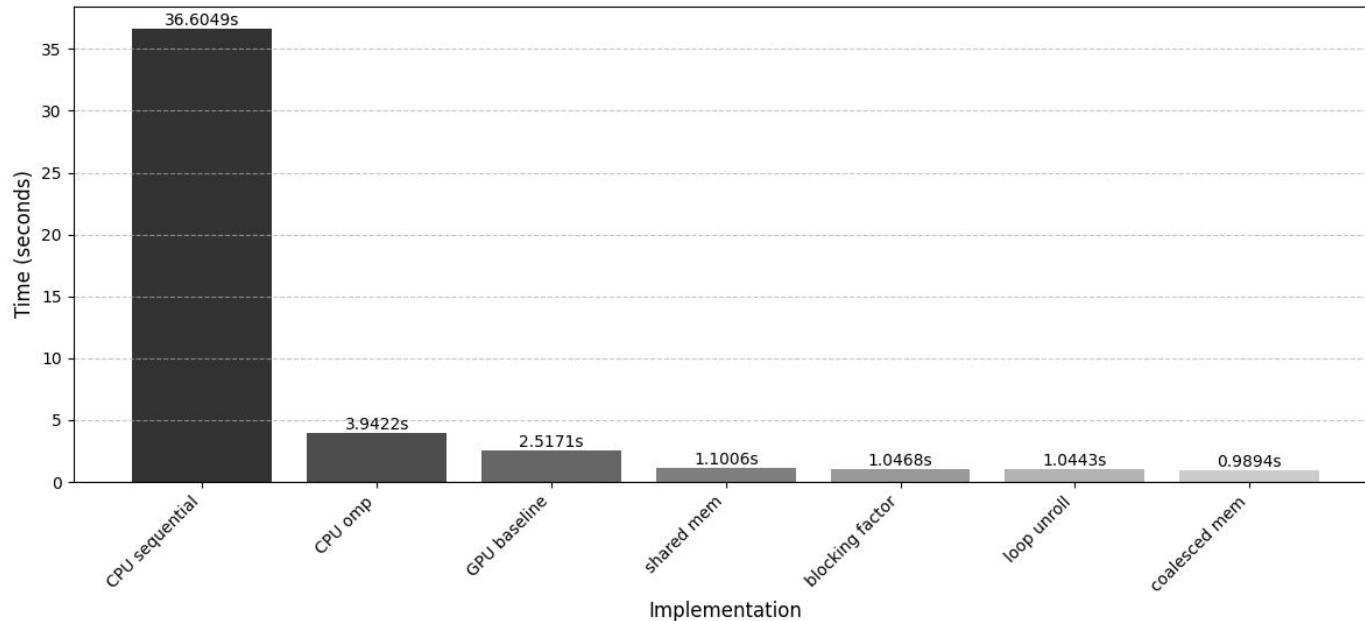
```
● pp24s057@apollo-login:~/filter$ srun -p nvidia -N1 -n1 --gres=gpu:1 nvprof --metrics sm_efficiency,shared_load_throughput,shared_store_th  
roughput,gld_throughput,gst_throughput ./filter flower.png out.png oil 5  
==1519514== NVPROF is profiling process 1519514, command: ./filter flower.png out.png oil 5  
==1519514== Some kernel(s) will be replayed on device 0 in order to collect all events/metrics.  
GPU time: 9.5478 seconds  
GPU time: 2.9458 seconds  
==1519514== Profiling application: ./filter flower.png out.png oil 5  
==1519514== Profiling result:  
==1519514== Metric result:  


| Invocations                                                                                 | Metric Name             | Metric Description             | Min        | Max        | Avg        |
|---------------------------------------------------------------------------------------------|-------------------------|--------------------------------|------------|------------|------------|
| Device "NVIDIA GeForce GTX 1080 (0)"                                                        |                         |                                |            |            |            |
| Kernel: oilPaintingKernelOptimized(unsigned char*, unsigned char*, int, int, int, int, int) |                         |                                |            |            |            |
| 1                                                                                           | sm_efficiency           | Multiprocessor Activity        | 99.97%     | 99.97%     | 99.97%     |
| 1                                                                                           | shared_load_throughput  | Shared Memory Load Throughput  | 25.299GB/s | 25.299GB/s | 25.299GB/s |
| 1                                                                                           | shared_store_throughput | Shared Memory Store Throughput | 562.00MB/s | 562.00MB/s | 562.00MB/s |
| 1                                                                                           | gld_throughput          | Global Load Throughput         | 420.20MB/s | 420.20MB/s | 420.20MB/s |
| 1                                                                                           | gst_throughput          | Global Store Throughput        | 160.57MB/s | 160.57MB/s | 160.57MB/s |
| Kernel: oilPaintingKernel(unsigned char*, unsigned char*, int, int, int, int, int)          |                         |                                |            |            |            |
| 1                                                                                           | shared_load_throughput  | Shared Memory Load Throughput  | 0.00000B/s | 0.00000B/s | 0.00000B/s |
| 1                                                                                           | shared_store_throughput | Shared Memory Store Throughput | 0.00000B/s | 0.00000B/s | 0.00000B/s |
| 1                                                                                           | gld_throughput          | Global Load Throughput         | 4.7946GB/s | 4.7946GB/s | 4.7946GB/s |
| 1                                                                                           | gst_throughput          | Global Store Throughput        | 21.285MB/s | 21.285MB/s | 21.285MB/s |

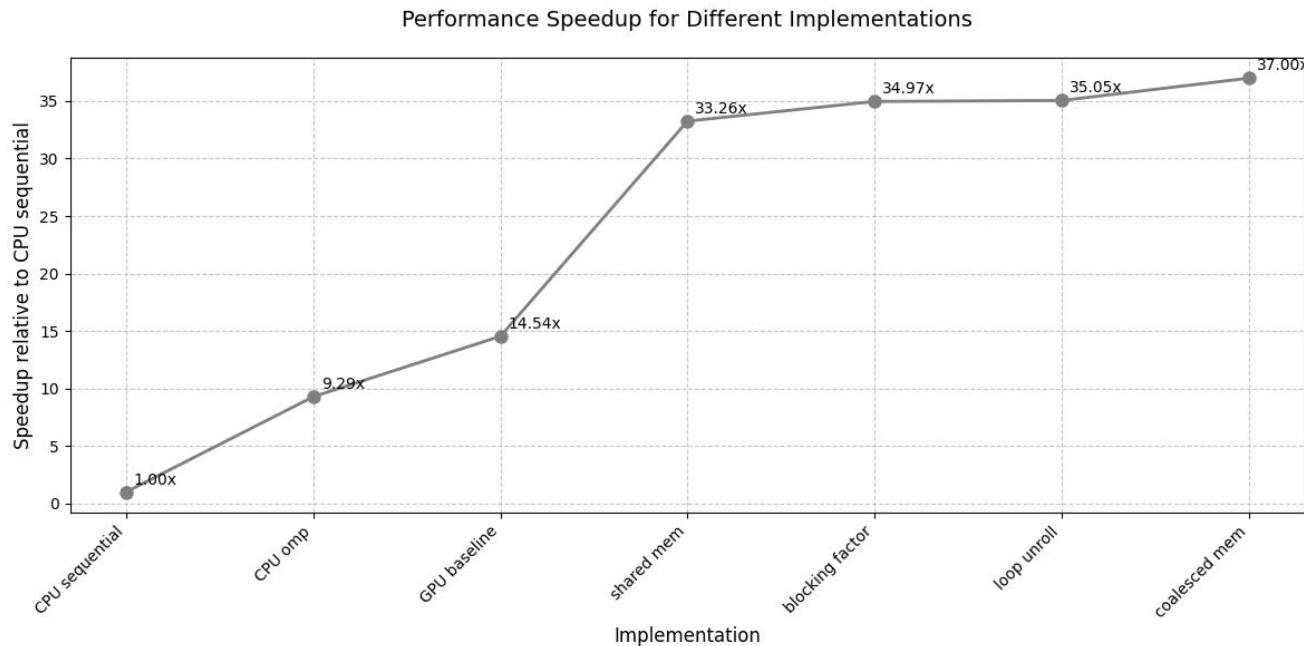

```

# CUDA: Baseline vs Optimized

Execution Time for Different Implementations

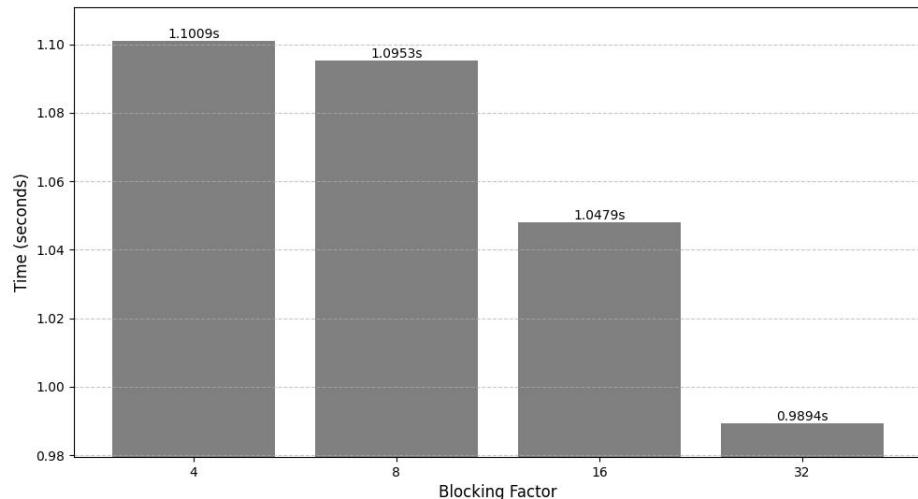


# CUDA: Baseline vs Optimized

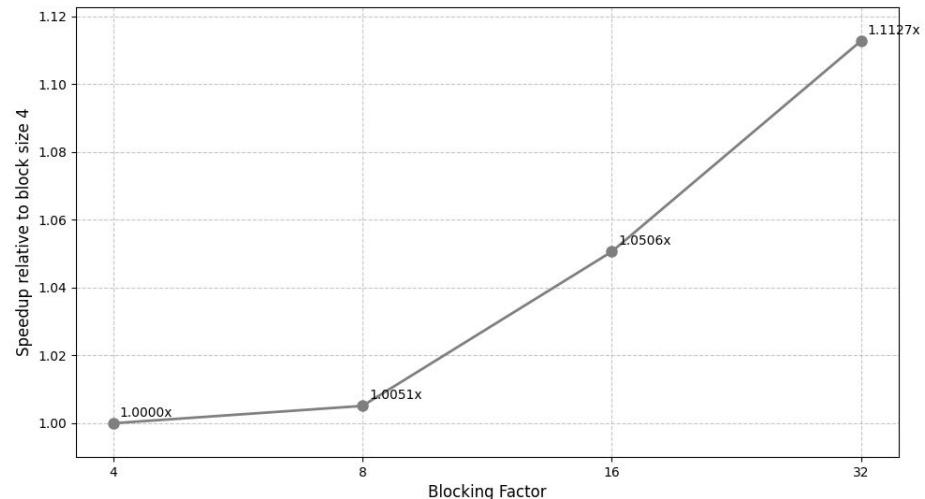


# CUDA: Blocking Factor

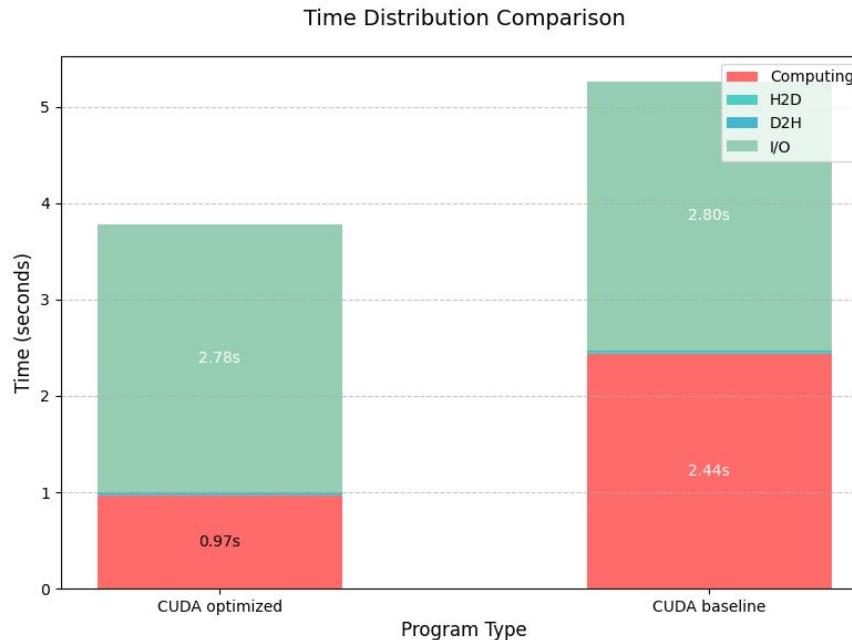
Execution Time vs Blocking Factor

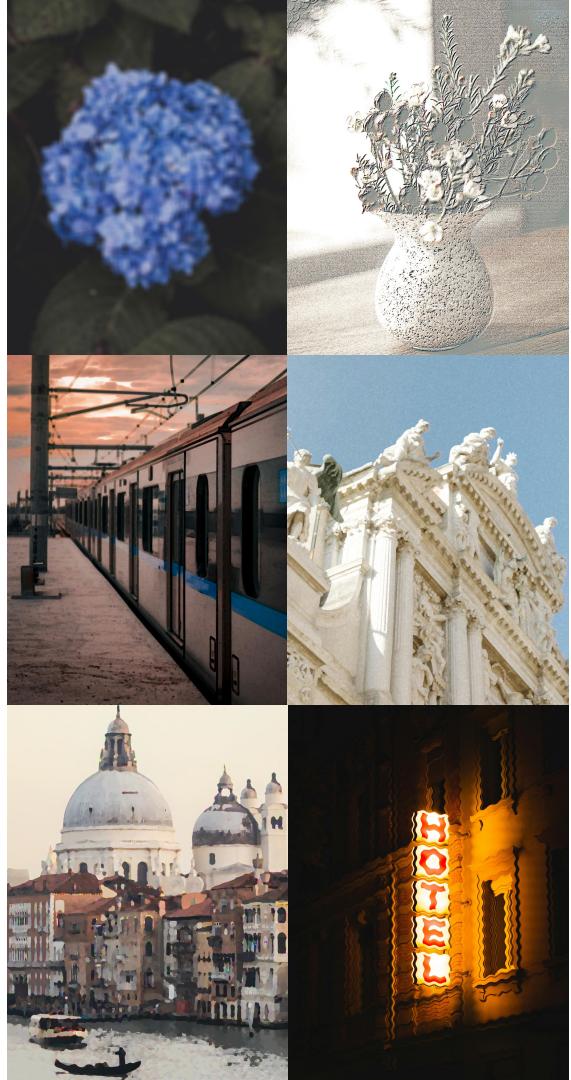


Performance Speedup vs Blocking Factor



# CUDA: Time Distribution





# Thank You

Team 8  
112062520 戴維恩  
111062698 戴樂為